

Socket programming

Complement for the programming assignment

INFO-0010

Outline

- Prerequisites
- Socket definition
- Briefing on the Socket API
- A simple example in Java
- Multi-threading and Synchronization
- Debugging tools
- HTTP Protocol
- Project overview

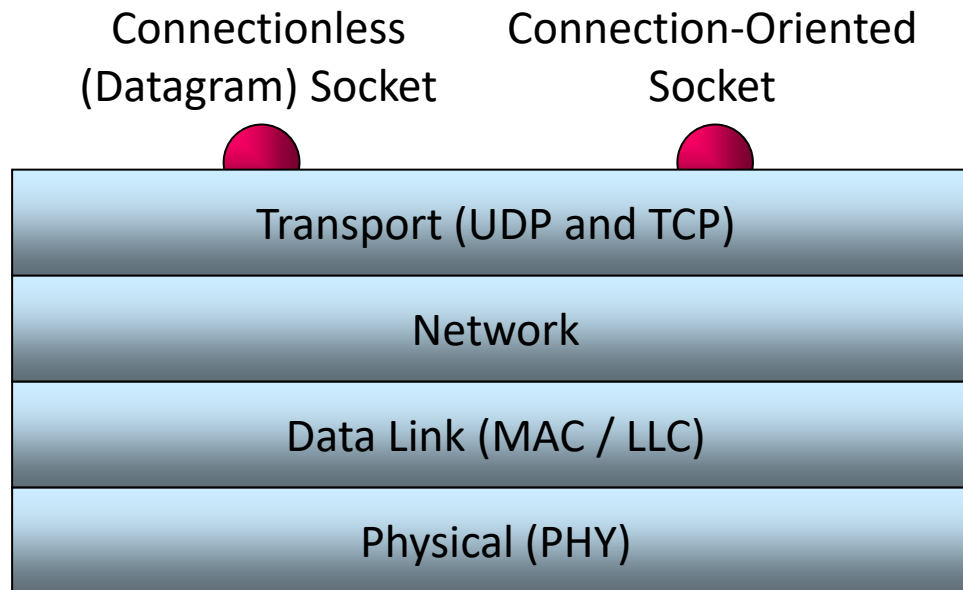
Prerequisites

Java Programming

- Compiling and executing Java programs
- Comments and Indentation
- Including libraries
- Classes, objects, methods, constructors
- Inheritance and Implementation
- Native types : boolean, byte, int, char, float, and basic operations on them.
- Simple arrays : [], ArrayList
- Alternatives and loops : if, while, for
- String manipulation : creation, concatenation, substring, comparison
- Thread creation
- Exceptions handling : try, catch, throws, throw.
- Input/OutputStreams usage : read, write, flush

What are sockets?

- Interface to network protocol stack
 - Typically to transport layer



What are sockets? (2)

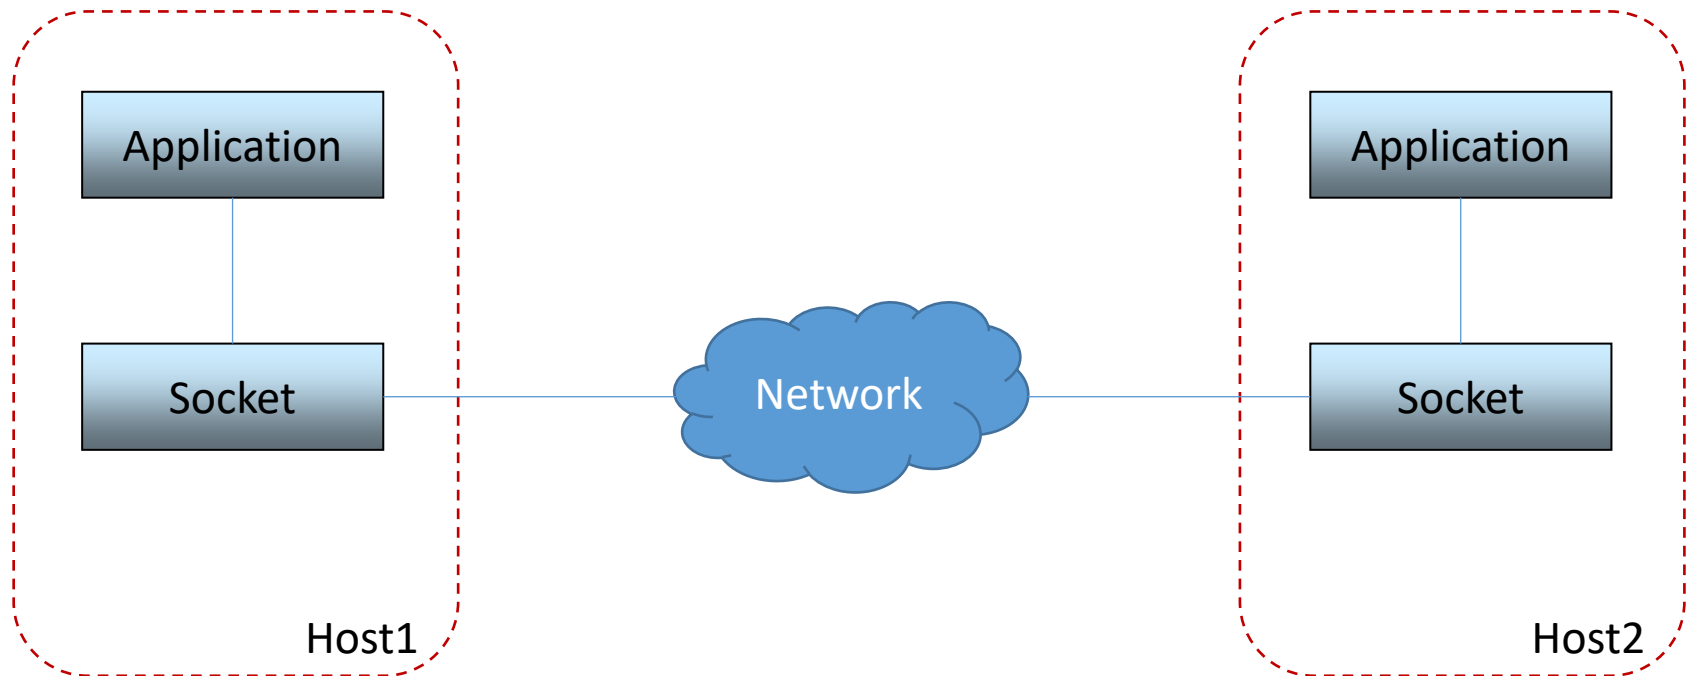
- A socket is an end-point of communication which identifies a local “process” at one end of a communication association
 - A socket is a half association
 - { protocol, local-address, local-port }
- A communication association is identified by two half associations
 - {
 protocol,
 local-address, local-port,
 remote-address, remote-port
}

Communication models

- Datagrams (UDP)
 - Message-oriented
 - Connectionless
 - Unreliable
 - No congestion control
- Connections (TCP)
 - Stream-oriented
 - Requires a connection
 - Reliable (no packet loss, no reordering)
 - Congestion control
- Raw
 - For traffic generation/capture

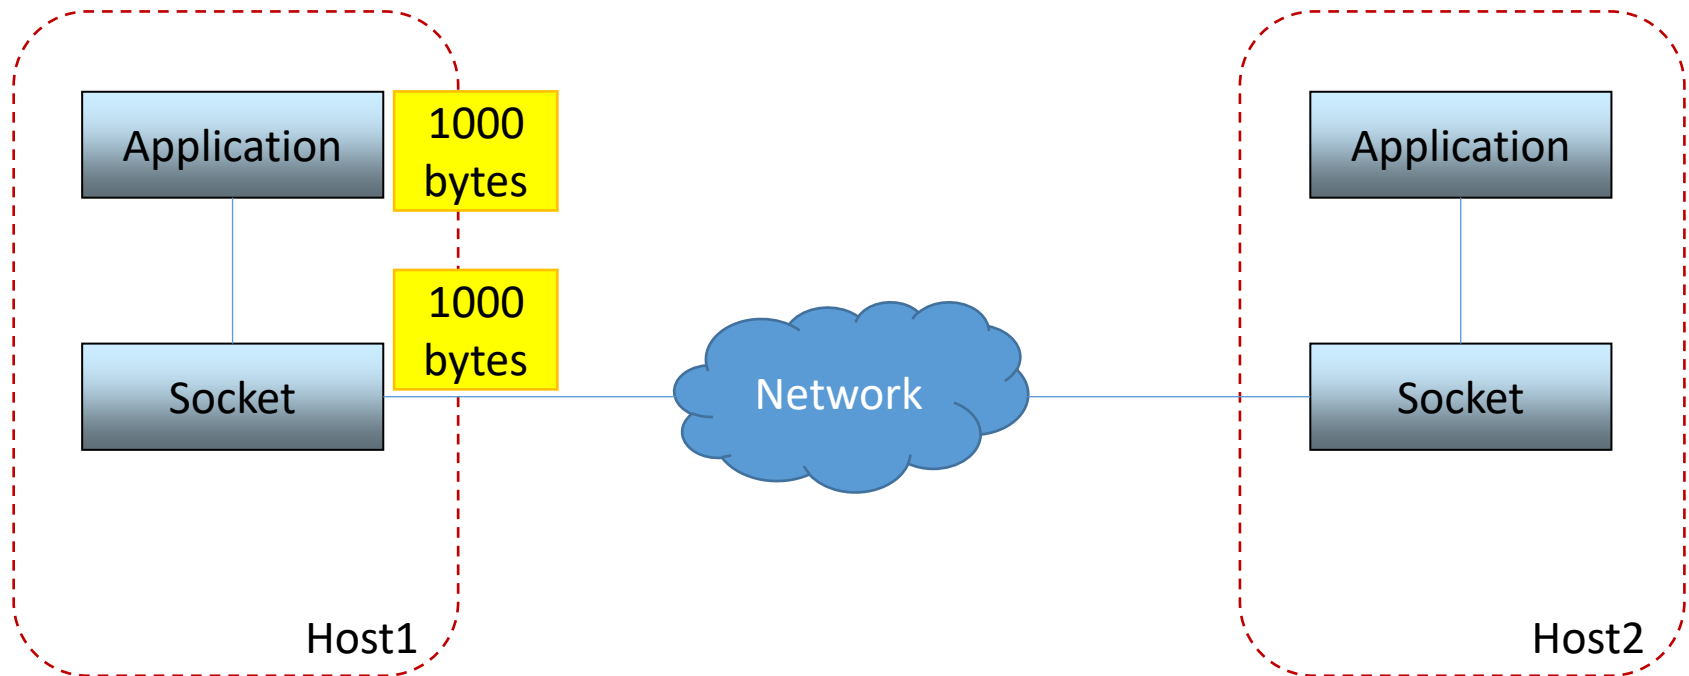
UDP vs TCP

- Conceptually:



UDP vs TCP

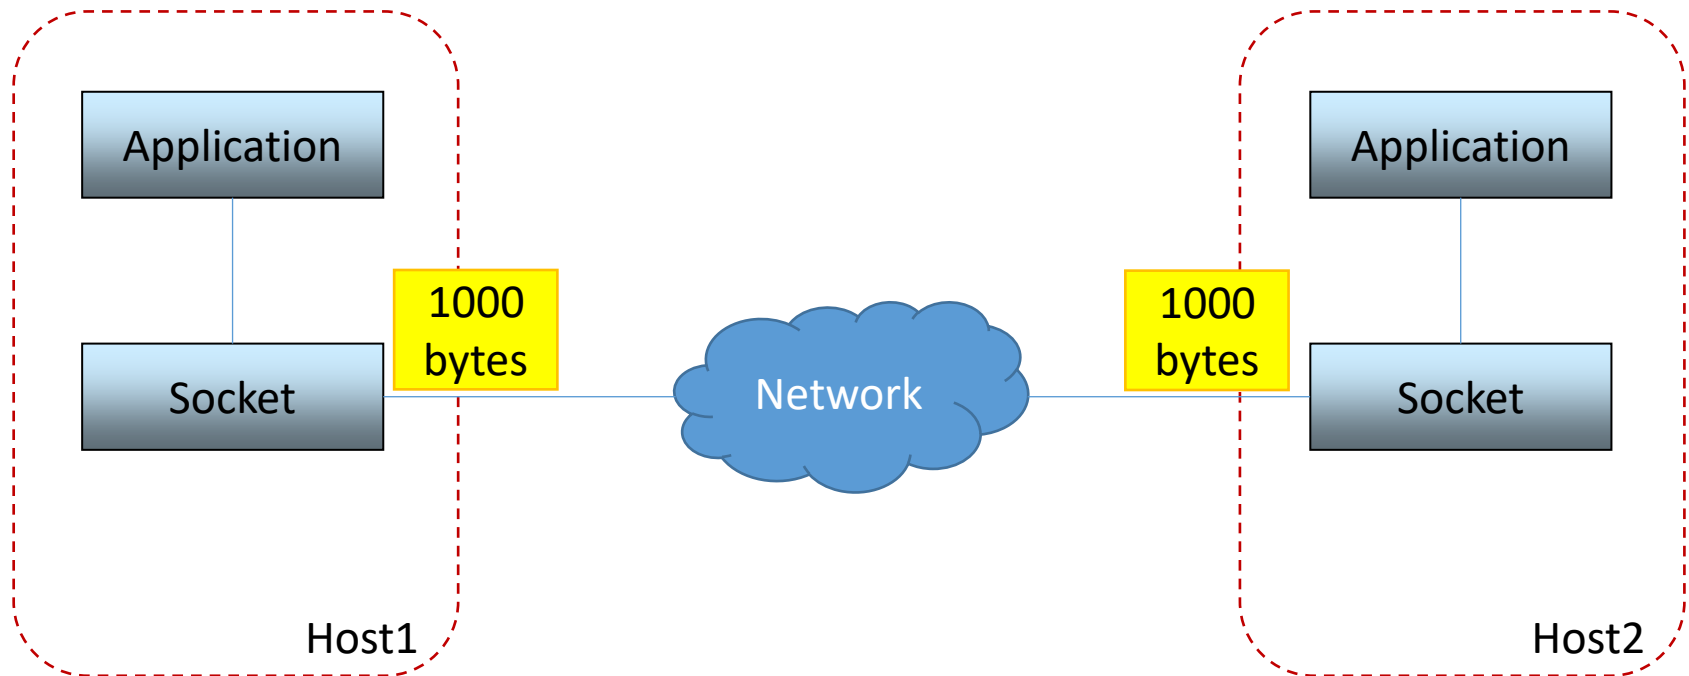
- Conceptually:



UDP

UDP vs TCP

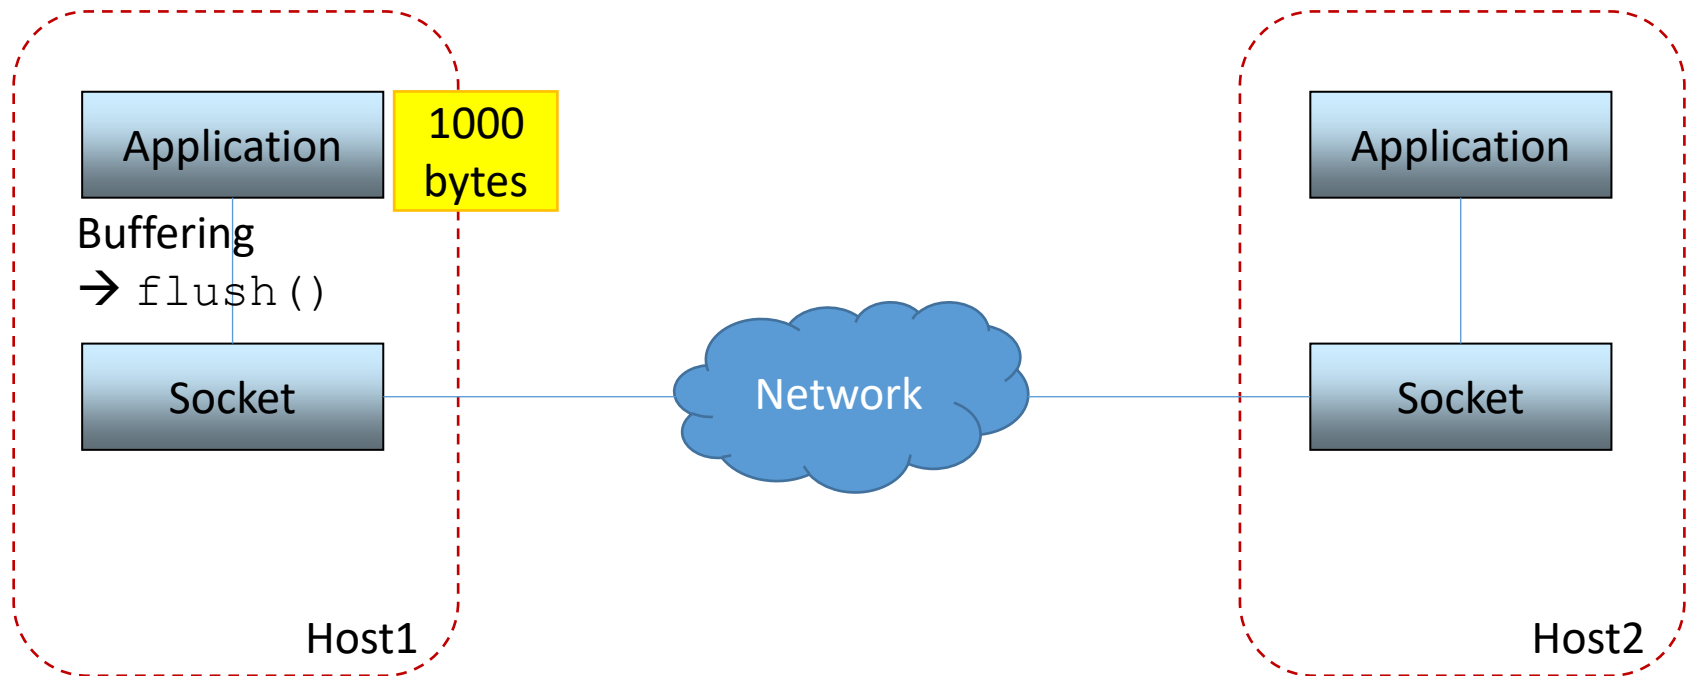
- Conceptually:



UDP

UDP vs TCP

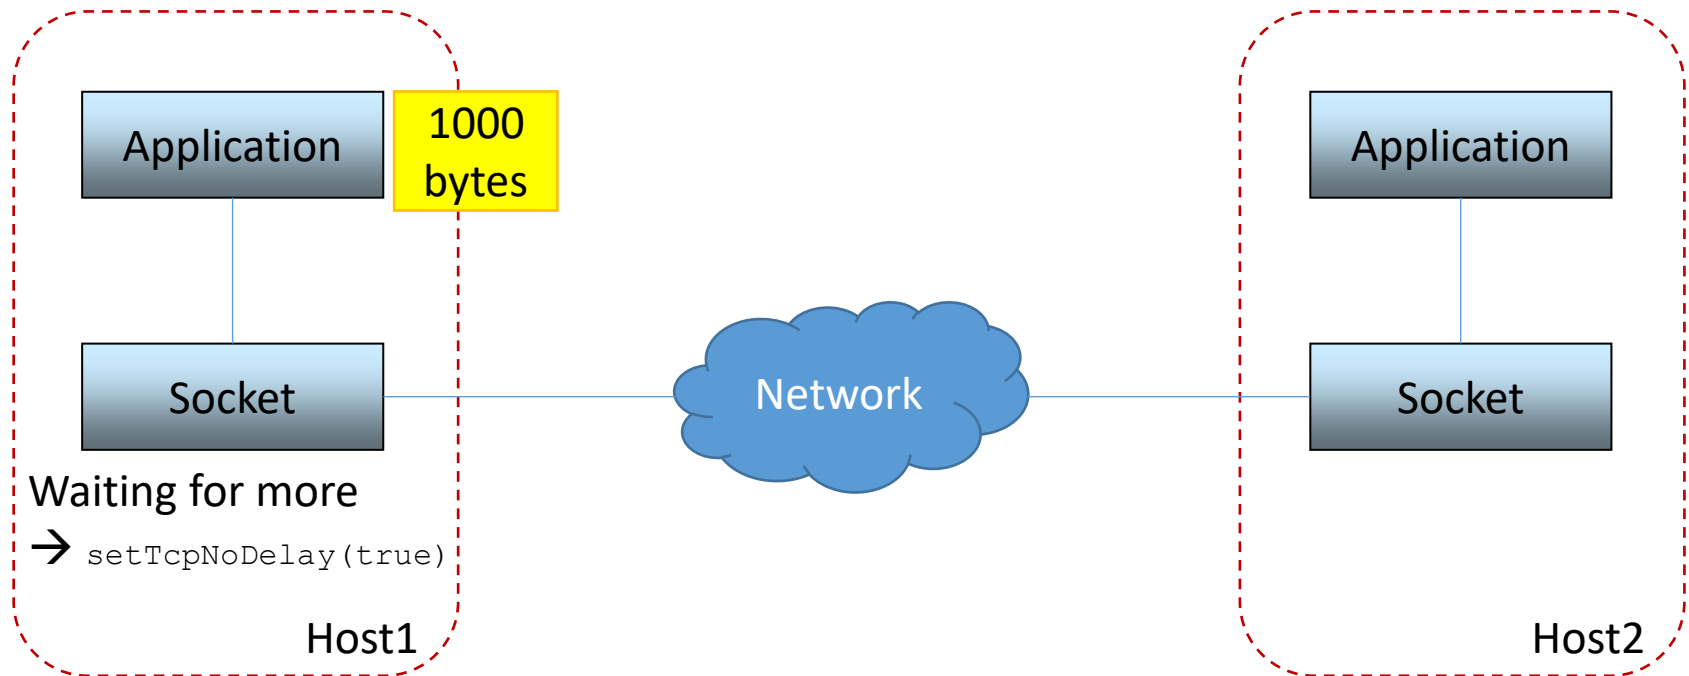
- Conceptually:



TCP

UDP vs TCP

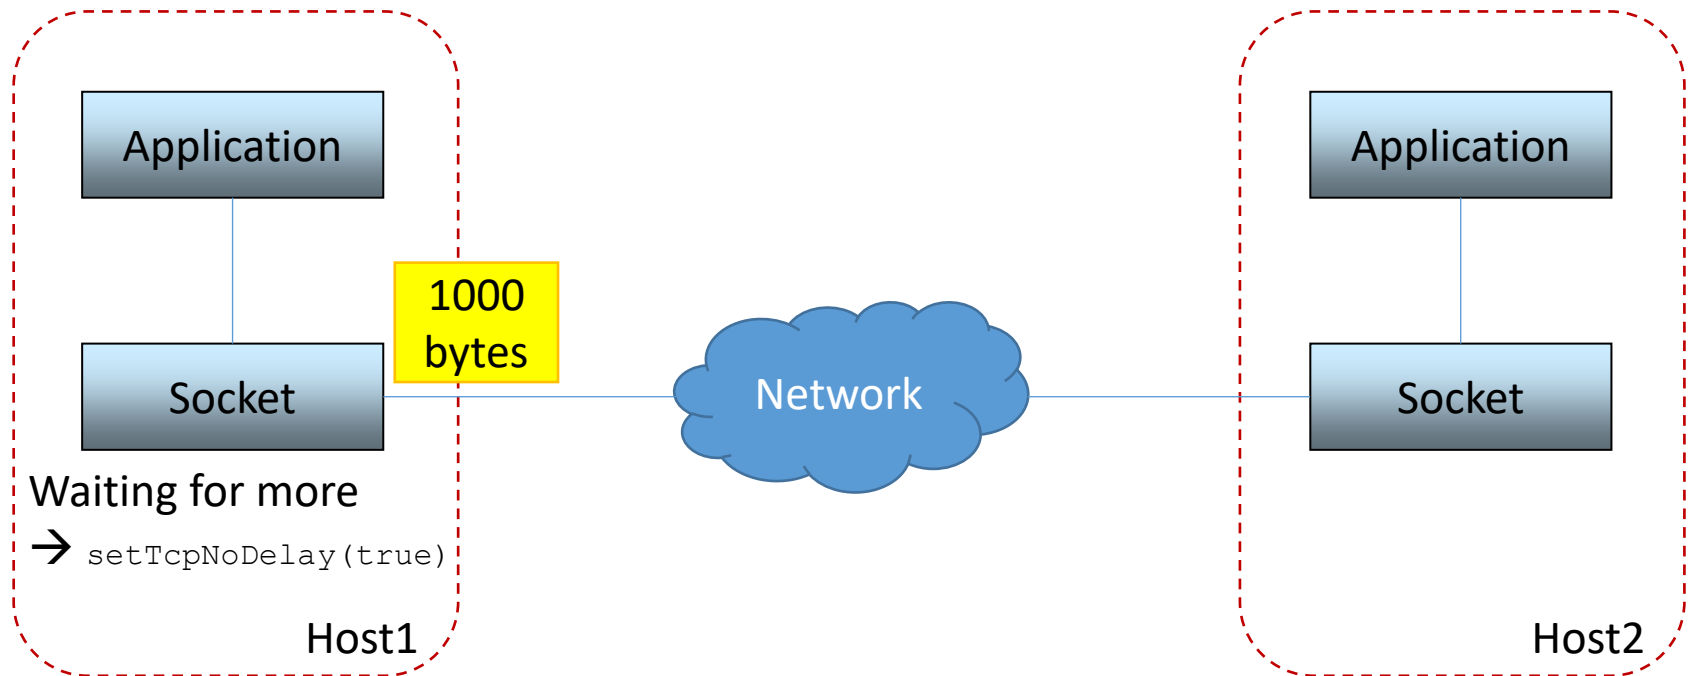
- Conceptually:



TCP

UDP vs TCP

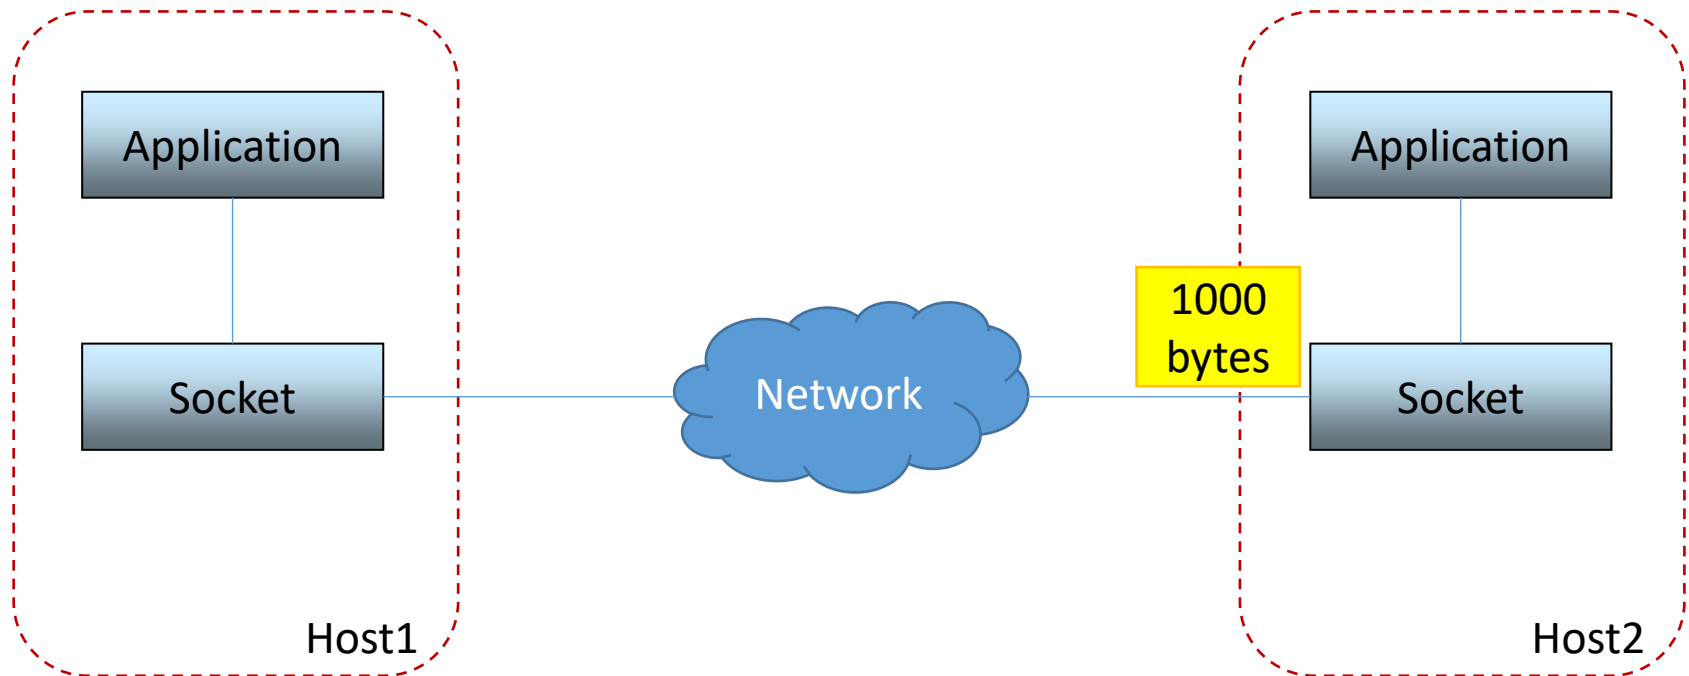
- Conceptually:



TCP

UDP vs TCP

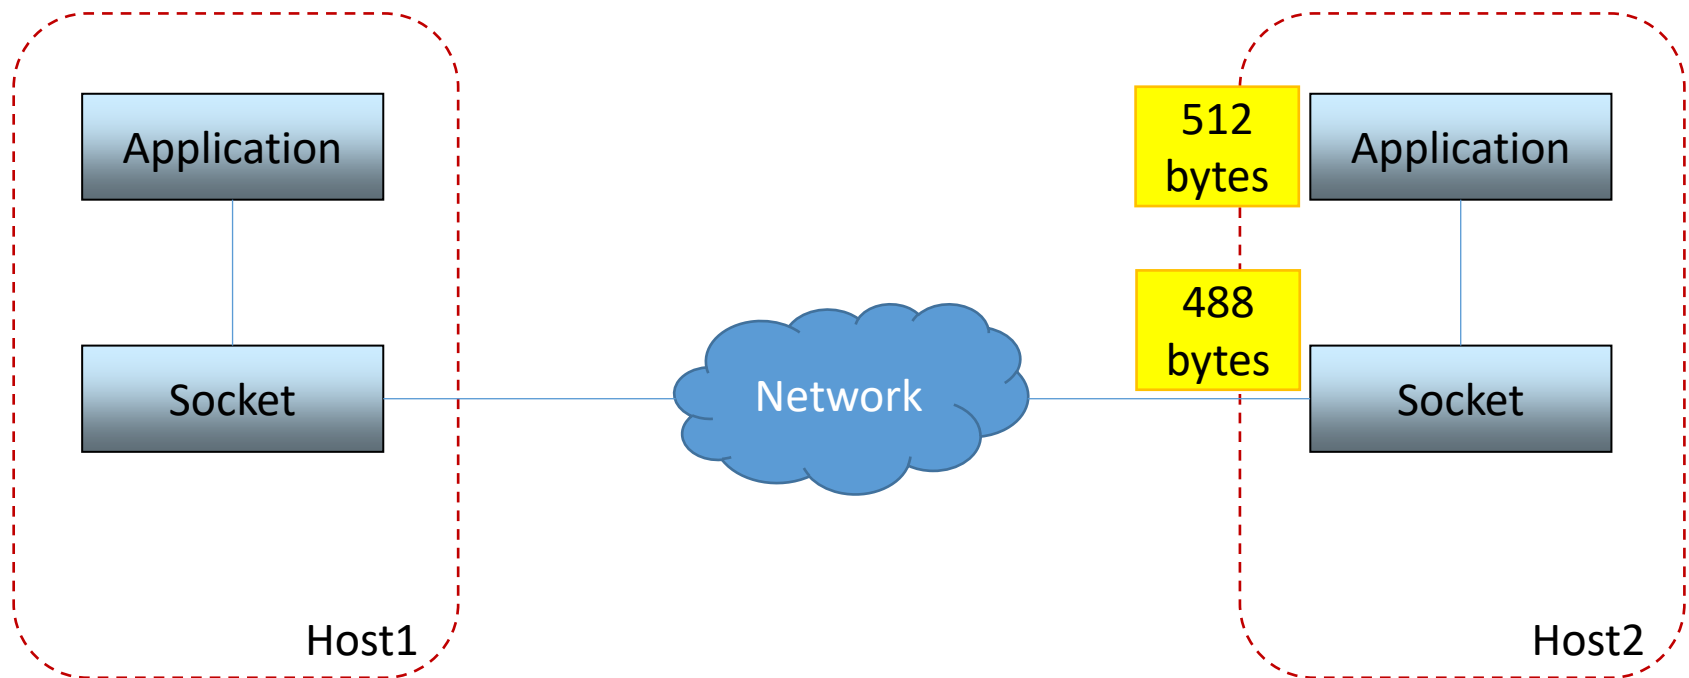
- Conceptually:



TCP

UDP vs TCP

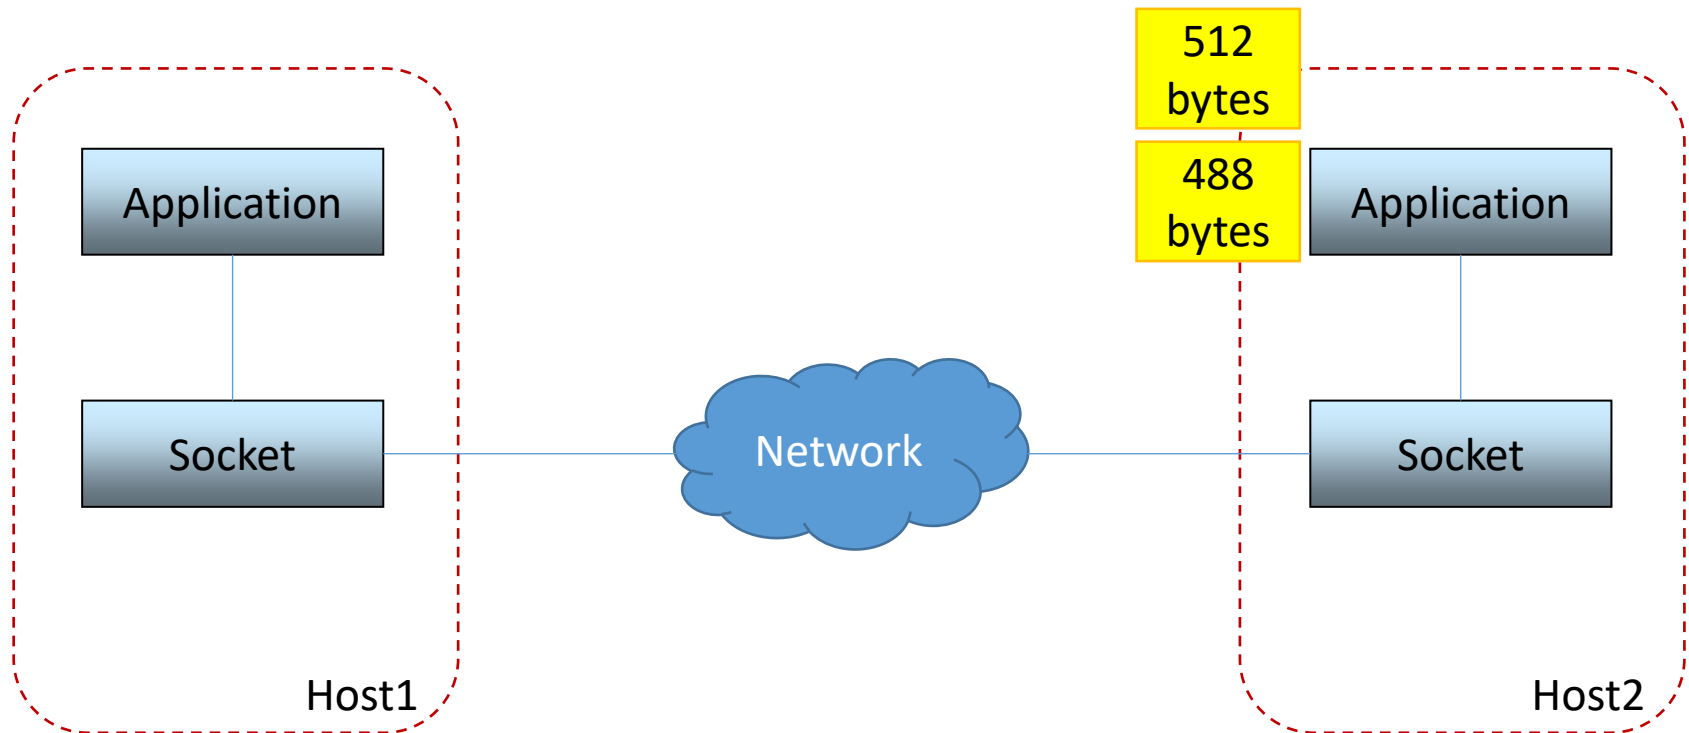
- Conceptually:



TCP

UDP vs TCP

- Conceptually:



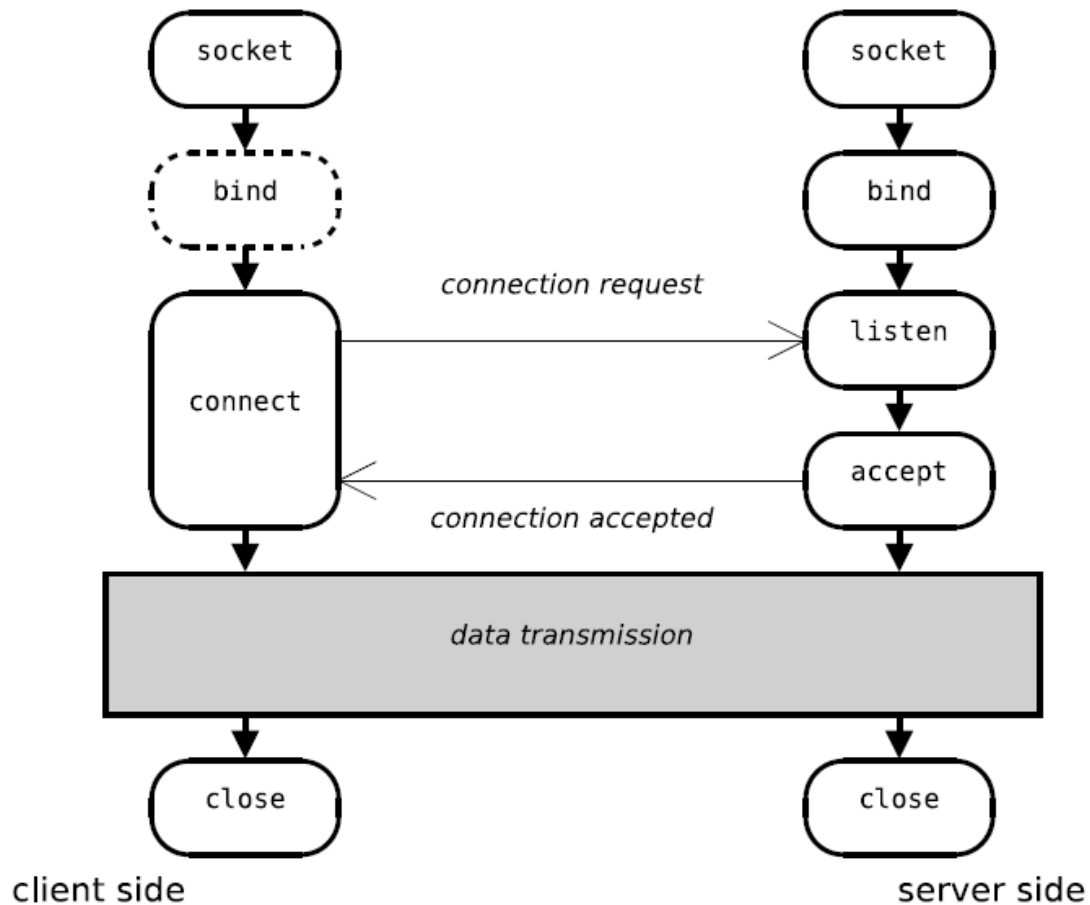
TCP

Connections

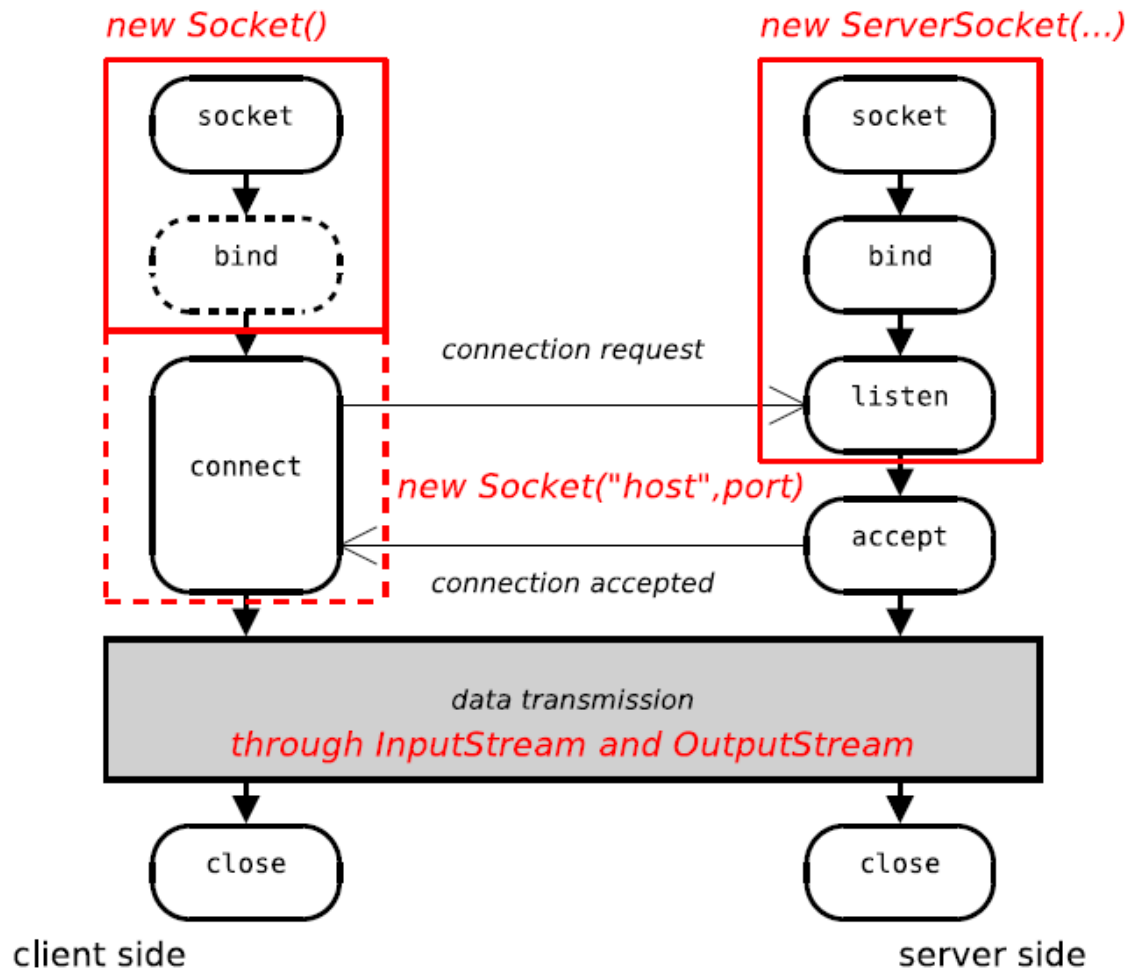
Implemented by TCP

- Reliable stream transfer
- Guarantees delivery and ordering provided connection not broken
- Does congestion control
 - What you have sent to the socket may not have left the box yet!
 - ✓ You can use `out.flush()` to force the writing to the socket
 - ✓ You can use `socket.setTcpNoDelay(true)` to disable Nagle's algorithm
- Chunks read may be different from chunks sent, but streams are identical
 - Programmer must check how many bytes should be read
 - Convention or application protocol header

Sockets' life cycle (syscalls)

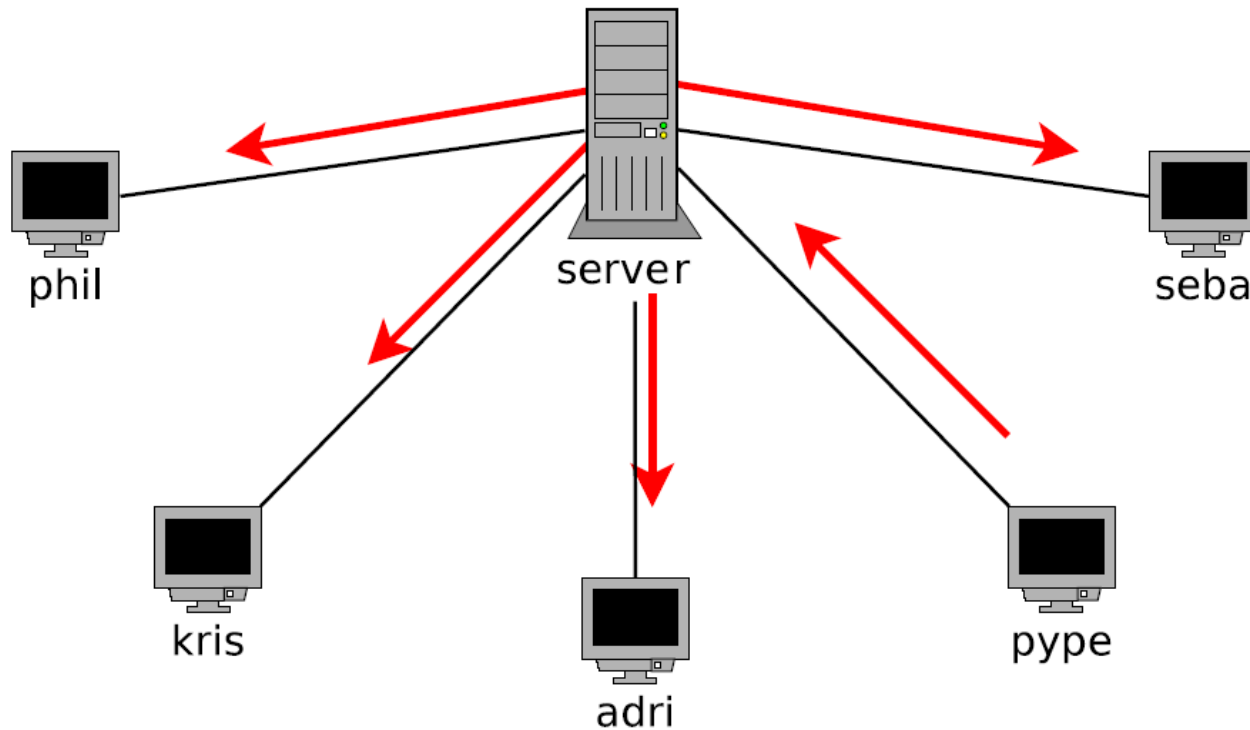


Socket's life cycle (Java)



Simple example : client/server chat

- Objective : The server duplicates each incoming message to all the clients.



- For starters, the server is only an echo server

The Client Side – a bot

```
import java.io.*;
import java.net.*;
```

```
class Bot {
```

```
    public static void main (String argv [ ] ) throws Exception {
```

```
        Socket s = new Socket ( "localhost", 8086 );
```

```
        OutputStream out = s.getOutputStream ( );
```

```
        InputStream in = s.getInputStream ( );
```

```
        byte msg[ ] = new byte [64];
```

```
        out.write ( "ANnA joined the channel" .getBytes ( ) );
```

```
        while ( true ) {
```

```
            if (in.read(msg) <= 0) break;
```

```
            if (new String(msg).startsWith("ANnA") )
```

```
                out.write( "ANnA feels fine , thanks.\n" . getBytes ( ) );
```

```
        }
```

```
        s.close ( );
```

```
    }
```

```
}
```

main() should always catch exceptions.
Would only make the code harder to read in this example.

Use explicit variable names.

We mainly focus on the class names in this example.

The Server Side – Incoming Connection

```
class Server {  
    public static void main ( String argv [ ] ) throws Exception {  
        ServerSocket ss = new ServerSocket (8086) ;  
        while ( true ) {  
            Socket ts = ss.accept () ;  
            OutputStream out = ts.getOutputStream() ;  
            InputStream in = ts.getInputStream() ;  
            out.write("Hello, this is the echo server".getBytes()) ;  
            byte msg [ ] = new byte [64] ;  
  
            while ( true ) {  
                int len= in.read(msg);    // get bytes (max 64)  
                if (len <=0) break ;        // connection closed by peer ?  
                out.write(msg,0,len);      // send them away .  
                out.flush();                // don't wait for more .  
            }  
            ts.close () ;  
        }  
    }  
}
```

What if multiple clients connect simultaneously?

The Server Side – multithreading

```
class Server {  
    public static void main ( String argv [ ] ) throws Exception {  
        ServerSocket ss = new ServerSocket (8086) ;  
        while ( true ) {  
            Socket ts = ss.accept () ;  
            Worker w = new Worker(ts);  
            w.start();           //Worker extends Thread  
        }  
    }  
}
```

- We spawn a *thread* every time a connection arrives
- That fresh-new thread will deal with the new client.
- And the main thread can return to the welcoming of incoming clients.

The Server Side – defining a thread

```
class Worker extends Thread {
    Socket s ;
    Worker ( Socket _s ) { s = _s ; }
    @Override
    public void run ( ) {
        try {
            OutputStream out = s.getOutputStream() ;
            InputStream in = s.getInputStream() ;
            out.write("Hello, this is the echo server". getBytes() ) ;
            byte msg [ ] = new byte [64] ;

            while ( true ) {
                int len= in.read(msg);           // get bytes (max 64)
                if (len <=0) break ;              // connection closed by peer ?
                out.write(msg,0,len);             // send them away .
                out.flush();                      // don't wait for more .
            }
            s .close ();                          //acknowledge end of connection
        } catch ( Exception any ) {
            System.err.println("worker died " + any) ;
        }
    }
}
```

Shared objects

- What if some objects need to be manipulated by different threads?
- For instance, we could keep a list (e.g. `ArrayList`) of all the `OutputStream` and turn the "echo" server into a real chat server.
- Multiple threads will use elements of the list simultaneously
 - Execution is **concurrent** and **non-atomic**
 - Consistency is thus not ensured
- Solution: **only one thread at a time** can use the elements in the list
 - Deem the sending phase a **critical section**
 - Implement **mutual exclusion** over critical section – i.e. prevent multiple threads from entering at once
 - **synchronized** keyword

The Server Side – shared objects

```
// all is just an ArrayList where each OutputStream is .add() ed
// out is the OutputStream corresponding to the Socket from which we are receiving
void send ( byte msg [ ] , int len ) throws Exception {
    synchronized ( all ) {
        for ( Enumeration e = all.elements ( ) ; e.hasMoreElements() ; ) {
            OutputStream o = (OutputStream ) e.nextElement() ;
            if ( o != out ) {
                o.write(msg,0, len ) ;    // send them away .
                o.flush();                // don't wait for more .
            }
        }
    }
}
```

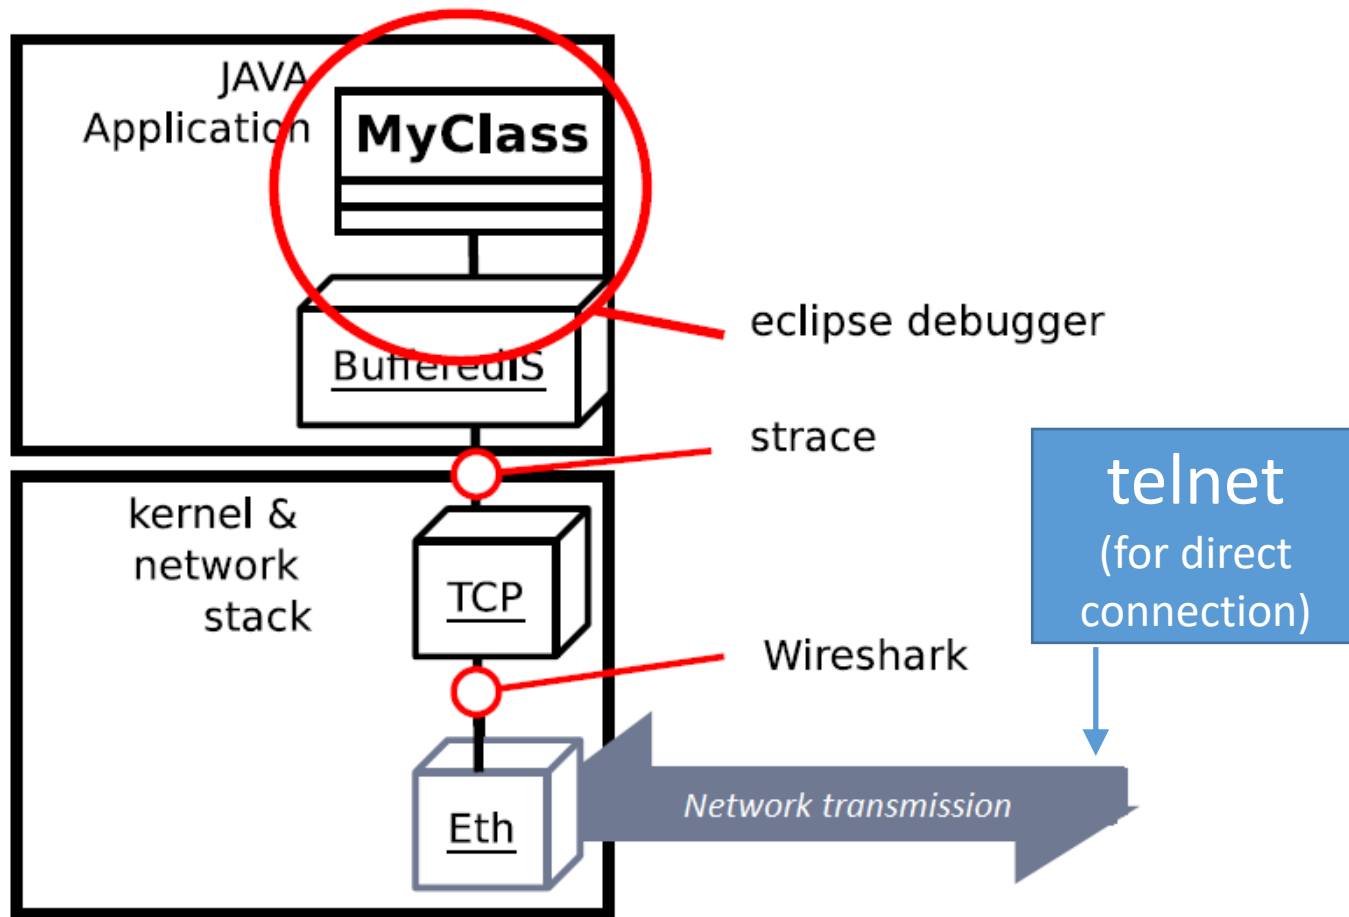
! Synchronization is done at the instance level !

For the same object, two threads cannot execute the critical section at the same time

Reading from TCP streams

- Recall : Chunks read may be different from chunks sent
- Sending a message on one side and calling **read** on the other side will not guarantee that the message has been fully read.
 - By the way, our chat bot is thus not very well coded.
- You must find the message boundaries
 - Message always have the same size X
 - While I haven't received X bytes, I keep calling **read** (possibly adjusting the number of bytes to read next)
 - Messages are separated by a delimiter
 - I keep reading until I see that delimiter
 - I might read more than one message, so the extra bytes must be considered as the beginning of a new message
 - If the delimiter is « `\r\n` », some classes might help (e.g. `BufferedReader`)
 - Messages are preceded with a header
 - The header has a fixed size
 - In the header, I find the size of the message.

A few debugging tools



Each bug has its proper catcher. So, use `catch` and `printStackTrace(...)` wisely!

strace -enetwork -f java Server

Pid

4199



```
[pid 4199] setsockopt(5, SOL_SOCKET, SO_REUSEADDR, [1], 4) = 0
[pid 4199] bind(5, {sa_family=AF_INET6, sin6_port=htons(8086), inet_pton(AF_INET6, ":::", &sin6_addr), sin6_flowinfo=0, sin6_scope_id=0}, 28) = 0
[pid 4199] listen(5, 50) = 0
[pid 4199] accept(5, {sa_family=AF_INET6, sin6_port=htons(3764), inet_pton(AF_INET6, "::ffff:127.0.0.1", &sin6_addr), sin6_flowinfo=0, sin6_scope_id=0}, [28]) = 6
Process 4211 attached
[pid 4211] sendto(-1251980320, umovestr: Input/output error 0xc, 3086176244, 0, ptrace: umoven: Input/output error {...}, 3042985144) = 0
[pid 4199] accept(5, <unfinished ...>)
[pid 4211] send(6, "Hello, this is the echo server", 30, 0) = 30
[pid 4211] recv(6, <unfinished ...>)
[pid 4199] <... accept resumed> {sa_family=AF_INET6, sin6_port=htons(3765), inet_pton(AF_INET6, "::ffff:127.0.0.1", &sin6_addr), sin6_flowinfo=0, sin6_scope_id=0}, [28]) = 7
Process 4215 attached
[pid 4215] sendto(-1252312096, umovestr: Input/output error 0xc, 3086176244, 0, ptrace: umoven: Input/output error {...}, 3042853368) = 0
[pid 4199] accept(5, <unfinished ...>)
[pid 4215] send(7, "Hello, this is the echo server", 30, 0) = 30
[pid 4215] recv(7, [
```

System calls

strace -enetwork -f java Server

Pid 4199

Pid 4211

Socket N° → 5 = new ServerSocket(8086)

```
[pid 4199] setsockopt(5, SOL_SOCKET, SO_REUSEADDR, [1], 4) = 0
[pid 4199] bind(5, {sa_family=AF_INET6, sin6_port=htons(8086), inet_pton(AF_INET6, ":::", &sin6_addr), sin6_flowinfo=0, sin6_scope_id=0}, 28) = 0
[pid 4199] listen(5, 50) = 0
[pid 4199] accept(5, {sa_family=AF_INET6, sin6_port=htons(3764), inet_pton(AF_INET6, "::ffff:127.0.0.1", &sin6_addr), sin6_flowinfo=0, sin6_scope_id=0}, [28]) = 6
Process 4211 attached
[pid 4211] sendto(-1251980320, umovestr: Input/output error 0xc, 3086176244, 0, ptrace: umoven: Input/output error {...}, 3042385144) = 0
[pid 4199] accept(5, <unfinished ...>)
[pid 4211] send(6, "Hello, this is the echo server", 30, 0) = 30
[pid 4211] recv(6, <unfinished ...>)
[pid 4199] <... accept resumed> {sa_family=AF_INET6, sin6_port=htons(3765), inet_pton(AF_INET6, "::ffff:127.0.0.1", &sin6_addr), sin6_flowinfo=0, sin6_scope_id=0}, [28]) = 7
Process 4215 attached
[pid 4215] sendto(-1252312096, umovestr: Input/output error 0xc, 3086176244, 0, ptrace: umoven: Input/output error {...}, 3042653368) = 0
[pid 4199] accept(5, <unfinished ...>)
[pid 4215] send(7, "Hello, this is the echo server", 30, 0) = 30
[pid 4215] recv(7, [
```

bytes to send

bytes handled by TCP

Wireshark

*Adapter for loopback traffic capture

Fichier Editer Vue Aller Capture Analyser Statistiques Telephonie Wireless Outils Aide

tcp.port == 2987 Filtering

No.	Time	Src Port	Dest Port	Protocol	Info
15	3.731806	49415	2987	TCP	49415 → 2987 [SYN] Seq=0 Win=65535 Len=0 MSS=65495 WS=256 SACK_PERM=1
16	3.731848	2987	49415	TCP	2987 → 49415 [SYN, ACK] Seq=0 Ack=1 Win=65535 Len=0 MSS=65495 WS=256 SACK_PERM=1
17	3.731905	49415	2987	TCP	49415 → 2987 [ACK] Seq=1 Ack=1 Win=2619648 Len=0
20	4.962343	49415	2987	TCP	49415 → 2987 [PSH, ACK] Seq=1 Ack=1 Win=2619648 Len=9
21	4.962422	2987	49415	TCP	2987 → 49415 [ACK] Seq=1 Ack=10 Win=2619648 Len=0
22	4.962626	2987	49415	TCP	2987 → 49415 [PSH, ACK] Seq=1 Ack=10 Win=2619648 Len=7
23	4.962652	49415	2987	TCP	49415 → 2987 [ACK] Seq=10 Ack=8 Win=2619648 Len=0
38	8.876625	49415	2987	TCP	49415 → 2987 [PSH, ACK] Seq=10 Ack=8 Win=2619648 Len=13
39	8.876682	2987	49415	TCP	2987 → 49415 [ACK] Seq=8 Ack=23 Win=2619648 Len=0
40	8.876932	2987	49415	TCP	2987 → 49415 [PSH, ACK] Seq=8 Ack=23 Win=2619648 Len=7
41	8.876960	49415	2987	TCP	49415 → 2987 [ACK] Seq=23 Ack=15 Win=2619648 Len=0
60	14.640570	49415	2987	TCP	49415 → 2987 [PSH, ACK] Seq=23 Ack=15 Win=2619648 Len=13
61	14.640625	2987	49415	TCP	2987 → 49415 [ACK] Seq=15 Ack=36 Win=2619648 Len=0
62	14.640976	2987	49415	TCP	2987 → 49415 [PSH, ACK] Seq=15 Ack=36 Win=2619648 Len=16
63	14.641004	49415	2987	TCP	49415 → 2987 [ACK] Seq=36 Ack=31 Win=2619648 Len=0
64	14.641045	2987	49415	TCP	2987 → 49415 [FIN, ACK] Seq=31 Ack=36 Win=2619648 Len=0
65	14.641061	49415	2987	TCP	49415 → 2987 [ACK] Seq=36 Ack=32 Win=2619648 Len=0
66	14.648315	49415	2987	TCP	49415 → 2987 [RST, ACK] Seq=36 Ack=32 Win=0 Len=0

< >

> Frame 38: 57 bytes on wire (456 bits), 57 bytes captured (456 bits) on interface \Device\NPF_Loopback, id 0

> Null/Loopback

> Internet Protocol Version 4, Src: 127.0.0.1, Dst: 127.0.0.1

> Transmission Control Protocol, Src Port: 49415, Dst Port: 2987, Seq: 10, Ack: 8, Len: 13

> Data (13 bytes)

```
0000 02 00 00 00 45 00 00 35 91 8d 40 00 80 06 00 00 .....E..5..@....
0010 7f 00 00 01 7f 00 00 01 c1 07 0b ab 0a 1c a8 f6 .....
0020 1e 75 e1 84 50 18 27 f9 4d de 00 00 54 52 59 20 ..u..P..'.M..TRY
0030 54 45 53 54 53 0d 0a 0d 0a TESTS.... Wordle protocol "TRY" request
```

wireshark_NPF_Loopback1P7EC2.pcapng

Paquets: 76 · Affichés: 18 (23.7%) · Perdus: 0 (0.0%) · Profile: Default

Some command lines

- (examples are better commented on the web).
- `javac Prog.java` to compile
- `java Prog` to launch
- `telnet localhost 8086` to test
- `strace -e trace=network -f java Prog` to track system calls issued by your program
- `netstat -tlp` to list server sockets
- `netstat -tcp` to list running connections.